

SMAI Assignment-4 Q1 Report

Name: Aditya Peketi
Roll: 2024122001

Question - 1

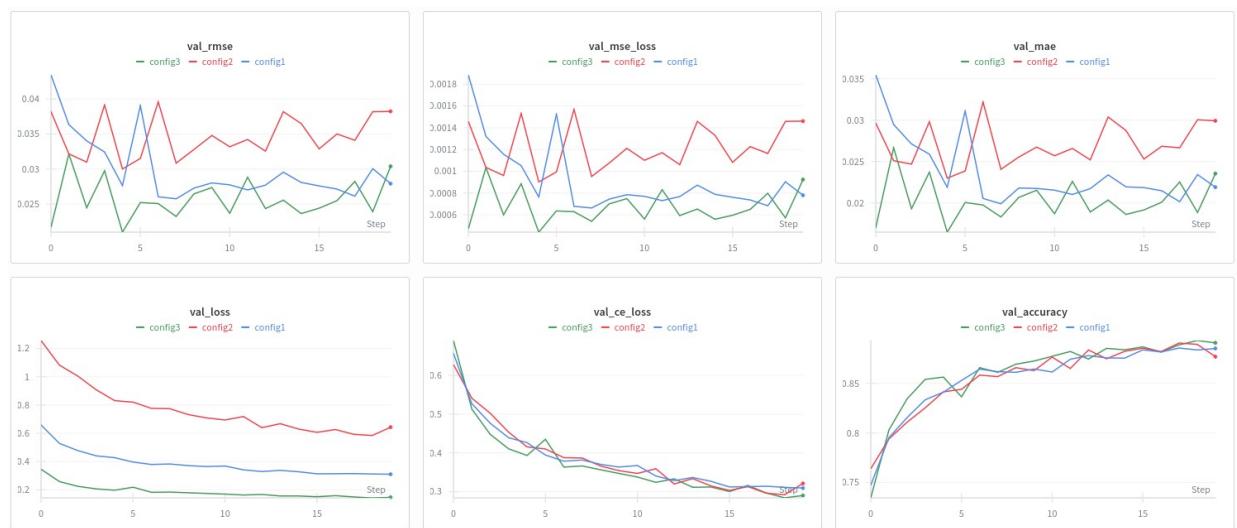
wandb Dashboard link: <https://api.wandb.ai/links/adityapeketii-iiit-hyderabad/5kg3ow7i>
workspace link - <https://wandb.ai/adityapeketii-iiit-hyderabad/SMAI-A4-multitask-fashion-mnist?nw=nwuseradityapeketii>

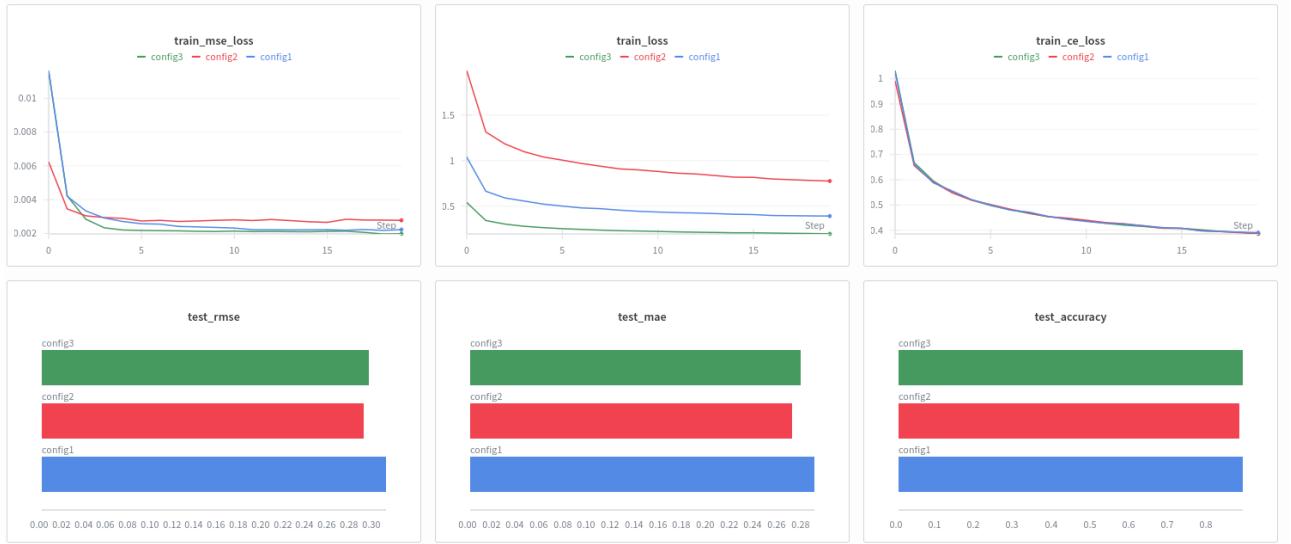
I have implemented a multi-task CNN with a shared convolutional backbone (2/3 Conv-BatchNorm-ReLU-Pool blocks) that branches into separate heads for classification (10 classes) and regression (ink prediction). The model optimizes a joint loss: $L = \lambda_1 LCE + \lambda_2 LMSE$.

I have tested and logged the model with **10 different configs**:

```
configs = [
    {'lambda1': 1.0, 'lambda2': 1.0, 'learning_rate': 0.001, 'dropout_rate': 0.3, 'optimizer': 'Adam', 'batch_size': 128, 'epochs': 20, 'num_layers': 3, 'base_filters': 32},
    {'lambda1': 2.0, 'lambda2': 0.5, 'learning_rate': 0.001, 'dropout_rate': 0.3, 'optimizer': 'Adam', 'batch_size': 128, 'epochs': 20, 'num_layers': 3, 'base_filters': 32},
    {'lambda1': 0.5, 'lambda2': 2.0, 'learning_rate': 0.001, 'dropout_rate': 0.3, 'optimizer': 'Adam', 'batch_size': 128, 'epochs': 20, 'num_layers': 3, 'base_filters': 32},
    {'lambda1': 1.0, 'lambda2': 1.0, 'learning_rate': 0.0008, 'dropout_rate': 0.3, 'optimizer': 'Adam', 'batch_size': 64, 'epochs': 20, 'num_layers': 2, 'base_filters': 16},
    {'lambda1': 1.0, 'lambda2': 1.0, 'learning_rate': 0.0008, 'dropout_rate': 0.3, 'optimizer': 'Adam', 'batch_size': 64, 'epochs': 20, 'num_layers': 2, 'base_filters': 32},
    {'lambda1': 1.0, 'lambda2': 1.0, 'learning_rate': 0.0005, 'dropout_rate': 0.3, 'optimizer': 'Adam', 'batch_size': 64, 'epochs': 20, 'num_layers': 3, 'base_filters': 16},
    {'lambda1': 2.0, 'lambda2': 1.0, 'learning_rate': 0.001, 'dropout_rate': 0.3, 'optimizer': 'Adam', 'batch_size': 128, 'epochs': 20, 'num_layers': 2, 'base_filters': 32},
    {'lambda1': 1.0, 'lambda2': 2.0, 'learning_rate': 0.0008, 'dropout_rate': 0.3, 'optimizer': 'AdamW', 'batch_size': 64, 'epochs': 20, 'num_layers': 2, 'base_filters': 32},
    {'lambda1': 0.5, 'lambda2': 1.0, 'learning_rate': 0.0005, 'dropout_rate': 0.25, 'optimizer': 'AdamW', 'batch_size': 64, 'epochs': 20, 'num_layers': 3, 'base_filters': 16},
    {'lambda1': 1.0, 'lambda2': 0.5, 'learning_rate': 0.002, 'dropout_rate': 0.35, 'optimizer': 'Adam', 'batch_size': 64, 'epochs': 20, 'num_layers': 3, 'base_filters': 16}
]
```

I wanted to test the trade off between classification and regression, thus in the first 3 runs I have only changed lambda1 and lambda2 keeping other parameters constant.





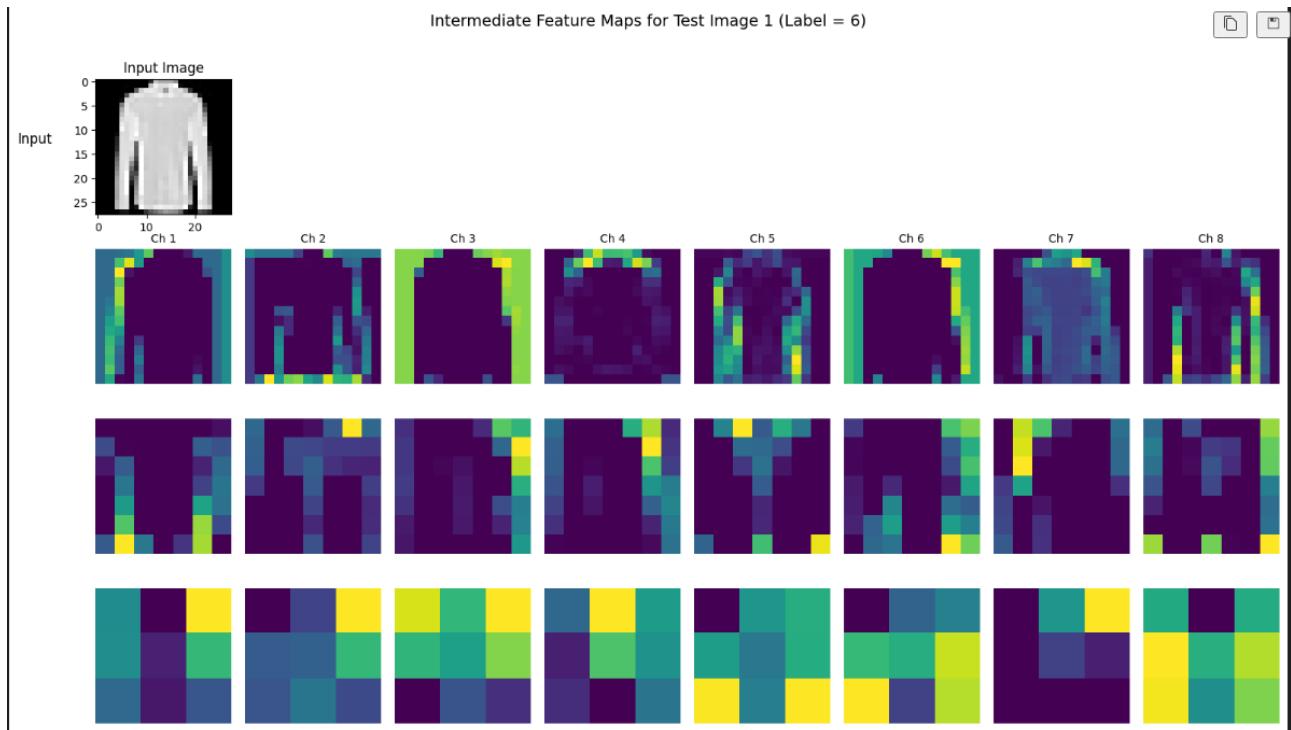
Observations:

When lambda1 is higher than lambda2 model prioritize classification over regression and this should result in a lower CE loss and higher classification accuracies.

And when lambda2 is higher than lambda1 model prioritize regression over classification and this results in a lower RMSE/ MSE loss and a lower classification accuracy

This can be seen from the above plots. The green represent Lambda2 > Lambda1 and red represents Lambda1 > Lambda2 and blue for both equal.

Feature Map Viz:



We can see that the first Conv 1 layer local edges of and borders of all the items. We can see bright regions around the contours of the clothing.

In conv layer 2 we can see more fine grain details like sleeves, bag edges and all.

All the metrics are in the wandb dashboard.

Best validation accuracy that I got is for config 3 and then config 1 respectively which are **89.083%** and **88.5%**. Other hyper parameters can be checked from the wandb link provided above.

Question- 2

Assumptions:

- Input Size: H x W (Default: 32x32, Doubled: 64x64)
- Convolutional Layers (Conv2d):
 - Kernel size: k x k
 - Stride: 1
 - Padding: k // 2
- Max Pooling Layers (MaxPool2d):
 - Kernel size: 2x2
 - Stride: 2 (This halves the spatial dimensions)
- Transposed Convolutional Layers (ConvTranspose2d):
 - Kernel size: 2x2
 - Stride: 2 (This doubles the spatial dimensions)
- Parameters: Biases and BatchNorm parameters are ignored in the calculations.
- Symbols:
 - NIC: Number of input channels
 - NF: Number of filters in the first convolutional layer
 - NC: Number of output color classes

1. Number of Weights

The number of weights in a convolutional layer is determined by:

`out_channels * in_channels * kernel_height * kernel_width`

Conv1 (Encoder): $NF * NIC * k * k$

Conv2 (Encoder): $(2NF) * NF * k * k$

Conv3 (Encoder): $(4NF) * (2NF) * k * k$

Deconv1 (Decoder): $(4NF) * (2NF) * 2 * 2$

Deconv2 (Decoder): $(2NF) * NF * 2 * 2$

Deconv3 (Decoder): $NF * NC * 2 * 2$

Classifier (1x1 Conv): $NC * NC * 1 * 1$

Total Weights (Symbolic Formula):

$$\text{Total_Weights} = (NF * NIC * k^2) + (2 * NF^2 * k^2) + (8 * NF^2 * k^2) + (8 * NF^2 * 4) + (2 * NF^2 * 4) + (NF * NC * 4) + NC^2$$

Simplified:

$$\text{Total_Weights} = k^2 * (NFCNIC + 10NF^2) + 40NF^2 + 4NF*NC + NC^2$$

This formula is valid for both 32x32 and 64x64 input sizes since weights do not depend on input resolution.

2. Number of Outputs (Activations)

Number of outputs = total count of activation elements across all layers for a single input image (channels * height * width).

Encoder1: $[NF, H/2, W/2] \rightarrow NF * (H/2) * (W/2)$

Encoder2: $[2NF, H/4, W/4] \rightarrow 2NF * (H/4) * (W/4)$

Encoder3: $[4NF, H/8, W/8] \rightarrow 4NF * (H/8) * (W/8)$

Decoder1: $[2NF, H/4, W/4] \rightarrow 2NF * (H/4) * (W/4)$

Decoder2: $[NF, H/2, W/2] \rightarrow NF * (H/2) * (W/2)$

Decoder3: $[NC, H, W] \rightarrow NC * H * W$

Classifier: $[NC, H, W] \rightarrow NC * H * W$

Total Outputs (Symbolic Formula):

$$\text{Total_Outputs} = (NFHW/4) + (2NFHW/16) + (4NFHW/64) + (2NFHW/16) + (NFHW/4) + (NCHW) + (NCH*W)$$

Simplified:

$$\text{Total_Outputs} = HW * [NF/4 + NF/8 + NF/16 + NF/8 + NF/4 + 2NC]$$

$$\text{Total_Outputs} = HW * [(4+2+1+2+4)NF/16 + 2NC]$$

$$\text{Total_Outputs} = HW * [13NF/16 + 2NC]$$

Calculations for Specific Input Sizes:

For 32x32 Input (H=32, W=32):

$$\text{Total_Outputs} = 1024 * [13NF/16 + 2NC] = 832NF + 2048NC$$

For 64x64 Input (H=64, W=64):

$$\text{Total_Outputs} = 4096 * [13NF/16 + 2NC] = 3328NF + 8192NC$$

(This is exactly 4 times the result for the 32x32 input.)

3. Number of Connections (MACs)

The number of connections (or Multiply-Accumulate operations) = (Total Weights in Layer) * (Output Spatial Area of that Layer).

For encoder Conv2d layers, computation happens before MaxPool2d.

Conv1: Weights = $NFNICk^2$, Output area = $HW \rightarrow Connections = (NFNICk^2)(HW)$

Conv2: Weights = $2NF^2k^2$, Output area = $(H/2)(W/2) \rightarrow Connections = (2NF^2k^2)(HW/4)$

Conv3: Weights = $8NF^2k^2$, Output area = $(H/4)(W/4) \rightarrow Connections = (8NF^2k^2)(HW/16)$

Deconv1: Weights = $8NF^2$, Output area = $(H/4)(W/4) \rightarrow Connections = (32NF^2)(HW/16)$

Deconv2: Weights = $2NF^2$, Output area = $(H/2)(W/2) \rightarrow Connections = (8NF^2)(HW/4)$

Deconv3: Weights = $4NFNC$, Output area = $HW \rightarrow Connections = (4NFNC)(HW)$

Classifier: Weights = NC^2 , Output area = $HW \rightarrow Connections = NC^2(H*W)$

Total Connections (Symbolic Formula):

$Total_Connections = HW * [NFNICk^2 + 2NF^2k^2/4 + 8NF^2k^2/16 + 32NF^2/16 + 8NF^2/4 + 4NFNC + NC^2]$

Simplified:

$Total_Connections = HW * [NFNICk^2 + 0.5NF^2k^2 + 0.5NF^2k^2 + 2NF^2 + 2NF^2 + 4NFNC + NC^2]$

$Total_Connections = HW * [k^2NF(NIC + NF) + 4NF^2 + 4NF*NC + NC^2]$

Calculations for Specific Input Sizes:

For 32x32 Input (H=32, W=32):

$Total_Connections = 1024 * [k^2NF(NIC + NF) + 4NF^2 + 4NF*NC + NC^2]$

For 64x64 Input (H=64, W=64):

$Total_Connections = 4096 * [k^2NF(NIC + NF) + 4NF^2 + 4NF*NC + NC^2]$

(This is exactly 4 times the result for the 32x32 input.)

Hyperparameter Tuning – Best Run

Validation Accuracy = 50.10645%

Validation Loss = 1.3

hyperarams - {'learning_rate': 1e-3, 'batch_size': 32, 'NF': 32, 'kernel_size': 5, 'optimizer': 'Adam'}

Link - <https://wandb.ai/adityapeketii-iiit-hyderabad/SMAI-A4-image-colorization>

Best model - <https://wandb.ai/adityapeketii-iiit-hyderabad/SMAI-A4-image-colorization/runs/9lcfmg4d>